

Bruno Masunaga

brunomasunaga@gmail.com | brunotatsuya.dev | linkedin.com/bruno-masunaga

Summary

Senior Backend Engineer with 7+ years building scalable, cloud-native distributed systems across fintech and platform products. Expertise in distributed systems, event-driven architecture, and observability, with proven experience in Go, Ruby, and Python. Focused on scalability, applied AI, agentic workflows, and developer productivity, delivering high-performance, resilient, and secure backend systems at scale.

Experience

Software Engineer, Qonto – France, Remote 2025 – present

- Designed and deployed an internal LLM microservice in Go using Kafka event-driven architecture and AWS Bedrock, applying model customization to fine-tune Llama models, achieving 95% accuracy in document data extraction and accelerating onboarding flows
- Built an agentic QA automation system based in LLM-driven agents (Claude Code) with learn behavior through path memoization and managed context, enabling reliable autonomous navigation to improve bug detection and quality assurance speed and precision
- Introduced AI-assisted development workflow based in spec-driven development (SDD) with context-managed coding agents to minimize iterations, significantly improving delivery speed by reaching ~80% right-first-time code generation
- Led adoption of an SLO framework across Go microservices, defining standards for SLIs/SLOs and alerting policies backed by Prometheus rules, reducing alert noise and improving incident detection lead time from hours to <1 min
- Designed and delivered a GraphQL gateway service using Go to decouple monolith into microservices without breaking clients, enabling incremental system decomposition
- Implemented Kafka Outbox pattern, ensuring reliable message delivery and reducing event-related production incidents by ~70%

Senior Software Engineer, Imaginary Cloud – Portugal, Remote 2024 – 2025

- Architected and led development of a cloud-native backend platform using Go, Kubernetes, Terraform, and GCP, designed to support ~20k concurrent users and ~1k requests/sec
- Designed and shipped multi-tenant architecture with strict tenant isolation and country-specific data residency enabling regulatory compliance
- Led a team of 4 cross-functional engineers, defining technical strategy and delivery standards that improved code review efficiency, QA throughput, and overall release velocity
- Conducted 15+ technical interviews across Junior to Lead levels, contributing to hiring 4 high-performing engineers and strengthening the company's delivery capabilities

Senior Software Engineer, AppTweak (via Imaginary Cloud) – Belgium, Remote 2024 – 2024

- Designed AI-powered backend services integrating Ruby async processing with LLM providers, deployed on AWS EKS using Terraform for scalable production workloads, achieving sustained throughput of ~30–100 messages/sec
- Re-architected request processing by offloading side effects to Sidekiq, improving Rails endpoint throughput, reducing user-facing timeouts and throttling from ~3% of requests to <0.5%
- Optimized data-intensive ETL pipelines by redesigning batch processing and tuning SQL queries, reducing AWS RDS burst balance consumption by 80% and improving long-term cost efficiency

Software Engineer, AppTweak (*via* Imaginary Cloud) – Belgium, Remote

2022 – 2023

- Led text search optimization by designing data indexing and replication strategies using Elasticsearch, reducing query latency from ~1 sec to ~100 ms per request (~90% faster) across Ruby on Rails service endpoints
- Redesigned Zendesk integration by implementing Redis-backed caching, reducing polling from 30 queries/hour to 1 query/hour and cutting database overhead by ~97%

Software Engineer, Itaú Unibanco – São Paulo, Brazil

2019 – 2022

- Led architecture and development of an internal low-code automation platform using Python and FastAPI, enabling 30+ non-technical users to create automated workflows
- Scaled data processing pipelines by redesigning Python and SQLAlchemy ETL workers with asynchronous integrations and multithreading, supporting 1,000+ daily workflows
- Provided technical leadership for a team of 5 engineers, guiding architecture decisions and delivery execution

Education

Bachelor of Science in Computer Science, Universidade Federal do ABC

Bachelor of Science in Science & Technology, Universidade Federal do ABC

Skills

Programming Languages: Go, Ruby, Python, TypeScript

AI & Agentic Systems: LLM application backends, AWS Bedrock, model customization, fine-tuning, agentic QA automation, specs-driven development (SDD), context-managed coding agents, workflow automation

Backend: Gin, GraphQL, gRPC, Kafka, Avro Schemas, Redis, PostgreSQL, MongoDB, OpenTelemetry, ELK (Elasticsearch, Logstash, Kibana), FastAPI, Ruby on Rails, Sidekiq, Prometheus, Grafana, Protobuf, Nginx

Cloud & Infrastructure: AWS (Bedrock, RDS, SQS, EKS, EC2, S3, IAM, SSM, MSK, ElastiCache), GCP (Cloud Run, Cloud Storage, Cloud SQL, GKE, Pub/Sub, IAM), Kubernetes, Terraform, Docker, Prometheus, CI/CD, ArgoCD

Core Strengths: Distributed systems, scalability, backend engineering, platform engineering, reliability engineering, observability, event-driven architecture, SLOs, microservices, AI-enabled developer tooling

Languages: English (Fluent), Portuguese (Native)